Engineering the Virtual Soundstage: Technical Analysis of Spatial Audio Systems and Software-Based Binaural Synthesis

The paradigm of audio reproduction has undergone a fundamental shift from channel-based layouts to object-based immersive environments. This transition, spearheaded by formats such as Dolby Atmos, Sony 360 Reality Audio, and Apple Spatial Audio, represents the culmination of decades of research into psychoacoustics, digital signal processing (DSP), and the physiological mechanics of human hearing.[1] In traditional stereophonic systems, audio is panned between two fixed loudspeakers, creating a limited "phantom image" primarily restricted to a horizontal plane between the transducers.[3] Spatial audio, by contrast, treats sound as discrete objects positioned within a virtualized 3D sphere, allowing for the perception of height, depth, and precise angular orientation.[5] For the consumer, the most accessible medium for this technology is binaural rendering over headphones, which attempts to trick the brain into perceiving an external soundstage through a standard two-channel output.[3] Replicating these effects in a software-only environment requires a sophisticated understanding of the Head-Related Transfer Function (HRTF), environmental acoustics, and the real-time constraints of mobile and desktop computing platforms.[10]

Core Principles of Spatial Audition and Psychoacoustic Mechanisms

The human ability to localize sound sources in three-dimensional space is a cognitive feat achieved through the interpretation of subtle temporal and spectral differences in the signals arriving at the eardrums.[1] This process, termed psychoacoustics, involves three primary vectors of data: horizontal localization (azimuth), vertical localization (elevation), and distance perception (depth).[1]

Horizontal Localization: ITD and ILD

Azimuthal perception relies on two main binaural cues: Interaural Time Difference (ITD) and Interaural Level Difference (ILD).[1] ITD occurs because a sound source located off-center will reach the ear closer to the source slightly earlier than the ear further away.[1] For low-frequency sounds (below 1.5 kHz), the brain utilizes phase differences and timing delays to calculate position.[1] As frequencies increase, the wavelength becomes smaller than the diameter of the

head, making phase-based timing cues ambiguous.[13]

At these higher frequencies, ILD becomes the dominant cue.[1] The human head acts as an acoustic shadow, attenuating high-frequency sound waves as they travel toward the "far" ear.[1] This attenuation, combined with the reflections from the torso and shoulders, provides a level difference that the brain interprets as a directional signal.[1]

Vertical Localization and the HRTF

Vertical localization and the resolution of the "cone of confusion"—where different points in space produce identical ITD and ILD values—are handled by spectral cues provided by the pinna (outer ear).[1] The folds and ridges of the pinna create a complex filter that introduces frequency-dependent peaks and notches into the incoming sound.[1]

Spectral Feature	Frequency Range	Perceptual Function
Primary Pinna Notch	6 kHz - 8 kHz	Core elevation cue; distinguishes height [13]
Secondary Pinna Notch	9 kHz - 12 kHz	Resolves front-back ambiguity [13]
Torso/Shoulder Reflection	200 Hz - 1.5 kHz	Provides scale and grounding for external sources [1]

The mathematical model of these filters is the Head-Related Transfer Function (HRTF), or its time-domain equivalent, the Head-Related Impulse Response (HRIR).[1] An HRTF is a pair of filters for every possible coordinate (θ, ϕ, r) around the listener.[12] When an audio signal is convolved with these filters, it is imprinted with the specific spatial signature the brain needs to place it in the environment.[9]

Distance and Depth Perception

Depth perception is a multi-modal process. In the free field, the primary cue is the inverse-square law, where the sound pressure level (SPL) decreases by 6 dB for every doubling of distance.[4] However, in enclosed environments, the ratio of direct-to-reverberant sound (D/R ratio) is the most critical cue.[2] As a source moves further from the listener, the intensity of the direct sound drops, while the intensity of the reverberant field (the sound reflected off room boundaries) remains relatively constant.[9] High-frequency air absorption also acts as a natural

low-pass filter, making distant sounds appear slightly "darker" than near sounds.[12]

Technical Architectures of Leading Spatial Audio Technologies

While the underlying psychoacoustic principles are universal, commercial implementations vary significantly in their delivery formats, metadata handling, and rendering pipelines.[5]

Dolby Atmos: The Hybrid Bed and Object Model

Dolby Atmos was originally designed for cinema and later adapted for home theater and music.[3] Its architecture is a hybrid of "beds" and "objects".[21] Beds are traditional channel-based buses (typically in 7.1.2 or 7.1.4 configurations) that provide a stable ambient foundation.[21] Objects are mono or stereo audio stems associated with metadata that defines their X, Y, and Z coordinates in three-dimensional space.[4]

The binaural rendering of Dolby Atmos for headphones utilizes a metadata-driven approach where creators assign a "Binaural Render Mode" to each object: Near, Mid, or Far.[3] These settings are measures of virtualized distance that dictate how the HRTF is applied.[21] Analysis suggests that Dolby Atmos binaural rendering is not a 100% HRTF convolution process; instead, it utilizes a blend of approximately 85% amplitude panning and 15% HRTF convolution.[14] This strategy mitigates the heavy spectral dips (10–15 dB) found in full HRTF sets, which can sometimes make music sound overly filtered or unnatural on certain headphones.[14]

Sony 360 Reality Audio: MPEG-H and Pure Objectivity

Sony 360 Reality Audio is based on the open MPEG-H 3D Audio standard (ISO/IEC 23008-3), developed in collaboration with Fraunhofer IIS.[5] Unlike Atmos, which often pre-renders objects into a binauralized stereo stream for certain consumer delivery paths, 360RA is a purely object-based format.[5] The device (smartphone, tablet, or AVR) receives up to 128 discrete audio objects and renders them locally.[25]

MPEG-H 3D Audio includes a sophisticated binaural renderer that employs a 64-band Complex Quadrature Mirror Filter (QMF).[28] This allows for frequency-dependent processing to mitigate the "comb filter" effect—an undesirable spectral artifact caused by phase differences between binaural filters when channel signals are downmixed.[28] In low-frequency bands, panning gains are compensated, while high-frequency bands undergo direct signal compensation to preserve timbral clarity.[28]

Apple Spatial Audio: Hardware-Software Co-Design

Apple Spatial Audio distinguishes itself through deep integration with the Apple ecosystem,

specifically leveraging the Inertial Measurement Units (IMUs) in AirPods and Beats headphones for dynamic head tracking.[29] The system is not a standalone format but a rendering layer that can ingest Dolby Atmos, multi-channel surround, or even upmix stereo content.[4]

Apple provides the Physical Audio Spatialization Engine (PHASE) for developers, which is a high-level API for simulating complex 3D acoustic environments.[33] PHASE handles not only positional audio but also geometry-based occlusion, where sound emitting from behind a virtual wall is attenuated based on the material properties of that wall.[33] The system also supports "Personalized Spatial Audio," where an iPhone's TrueDepth camera is used to create a 3D scan of the user's ears, allowing the system to customize the HRTF filters to the user's unique anatomy.[22]

Technology	Core Standard	Delivery Mechanism	Head Tracking
Dolby Atmos	ADM-BWF [14]	Hybrid Bed/Object [23]	Optional / Platform [4]
Sony 360RA	MPEG-H [5]	Pure Object-Based [5]	Device-specific [36]
Apple Spatial	PHASE / Core Audio	Rendering Layer [34]	System-level IMU [37]

Software replication of Spatial Audio: The DSP Pipeline

Recreating a high-fidelity spatial audio experience in software without specialized hardware requires a modular DSP engine that processes incoming audio objects through several stages of virtualization.[10]

The Binaural Rendering Core

At the heart of any spatial audio software is the binaural renderer.[8] For every audio object, the engine must perform a real-time convolution with a pair of HRTF filters corresponding to the object's current azimuth and elevation relative to the listener's head pose.[10] Because HRTF datasets are measured at discrete points (e.g., every 5 degrees), the engine must perform continuous interpolation as objects move.[15]

Linear interpolation of HRIRs in the time domain can cause "comb filtering" or phase cancellations.[10] Advanced systems often perform interpolation in the frequency domain or use Spherical Harmonics (SH).[13] By representing the HRTF as a series of SH coefficients, the renderer can achieve smooth, artifact-free transitions as a sound source or the listener's head

rotates.[13]

FFT-Based Convolution and Overlap-Add

For low-latency performance, convolution is typically performed in the frequency domain using the Fast Fourier Transform (FFT).[39] The standard method is the Overlap-Add (OLA) or Overlap-Save (OLS) algorithm, which segments the audio into blocks, performs FFT, multiplies the spectra by the HRTF, and performs the Inverse FFT (IFFT).[39] To maintain zero or ultra-low latency, the system may use partitioned convolution, where the impulse response is broken into small, non-uniform segments to minimize the delay inherent in larger FFT windows.[42]

Room Simulation and Environmental Modeling

A dry HRTF convolution often sounds "inside the head" because it lacks the acoustic context of a physical space.[3] To achieve realistic "out-of-head" localization, the software must simulate early reflections and late reverberation.[9]

The Image Source Method (ISM) is the standard geometric approach for modeling early reflections.[16] The engine calculates virtual "image sources" by mirroring the emitter across the room's surfaces.[16] Each reflection is treated as a delayed, attenuated, and filtered version of the original source, which is then binauralized based on its relative angle.[16]

For late reverberation, which is the dense, diffuse tail of the room response, Feedback Delay Networks (FDNs) are the preferred architecture.[18] FDNs use parallel delay lines connected via a unitary feedback matrix (such as a Householder or Hadamard matrix) to generate exponentially decaying noise that mimics the recursive sound reflections in a 3D volume.[18]

Room Acoustic Component	**Modeling Algorithm**	**Computational Complexity**
Direct Path	HRTF Convolution [12]	High (per source)
Early Reflections	Image Source Method [16]	Moderate (order dependent)
Late Reverb	Feedback Delay Network [18]	Low (global)

Cross-Device Compatibility and Personalization Strategies

The primary challenge in deploying spatial audio to "any headphones" is the variance in human anatomy.[1] A "generic" HRTF, often measured from a manikin like the KEMAR, might work well for 60% of the population but lead to significant localization errors for the remaining 40%.[43]

Generic vs. Personalized HRTF

Using a non-personalized HRTF can lead to "front-back confusion," where the listener cannot distinguish if a sound is coming from 0° or 180°.[1] It also impacts externalization—the sense that a sound is truly in the room rather than just at the ear entrance.[43]

To address this without specialized acoustic measurement labs, software developers are turning to machine learning for "HRTF Individualization".[13] Recent research has demonstrated that deep neural networks, specifically Transformers and GANs, can predict a personalized HRTF from a simple set of anthropometric measurements (e.g., ear height, head width) or a 2D photograph.[13] These models are trained on large public datasets like the SONICOM or CIPIC databases.[50]

Calibration and User Tuning

An alternative to full personalization is user-driven calibration.[12] Software can offer the following "tuning" parameters to help users align the virtual soundstage with their perception:

- **Virtual Distance Blend:** Adjusting the ratio of HRTF convolution to amplitude panning (similar to Dolby's 15/85 blend).[14]
- **Elevation Tilt:** Shifting the center frequency of the pinna notches to compensate for different ear shapes.[13]
- **Interaural Distance (ITD) Scaling:** Adjusting the timing delay to match the listener's head size.[1]

Performance and Real-Time Optimization

Spatial audio is computationally demanding. Rendering 128 objects with individual HRTFs, reflections, and reverb can easily overwhelm a mobile processor.[10]

SIMD and Parallel Processing

To achieve real-time performance, the rendering engine must be optimized for Single Instruction, Multiple Data (SIMD) instruction sets, such as Intel SSE/AVX or ARM NEON.[58] SIMD allows the processor to perform the same operation (e.g., gain multiplication or FFT butterfly) on multiple samples simultaneously.[58] For a stereo signal, de-interleaving the channels into separate SIMD registers during processing can provide a 2x to 5x speedup compared to

standard scalar code.[58]

Latency Constraints and Head Tracking

Latency is the "killer" of immersive experiences.[42] In head-tracked systems, the delay from head motion to audio update (motion-to-photon/motion-to-audio) must be below 20ms to prevent disorientation and "swimming" soundfields.[38] This requires extremely efficient sensor fusion of the IMU data and predictive filtering (such as a Kalman filter or linear extrapolation) to estimate the head position at the exact moment the next audio buffer will be heard.[62]

On Android, developers utilize the Oboe or AAudio APIs to request "Low Latency" and "Exclusive" modes, which bypass much of the OS's audio mixer to achieve sub-10ms round-trip latency.[64]

Platform	Optimization API	Typical Latency Target
Android	Oboe / AAudio [68]	< 15 ms
iOS / macOS	Core Audio / PHASE [34]	< 10 ms
PC / Gaming	ASIO / Wwise / FMOD [57]	< 20 ms

Conceptual Architecture Design of a Software Spatial Engine

A robust, cross-platform spatial audio engine is typically designed with a modular architecture that separates state management from audio rendering.[10]

The Input and Scene Management Layer

This layer ingests audio streams (mono, stereo, or multi-channel) and manages a "Scene Graph" of the 3D environment.[33] Every object in the scene has a state structure containing its position vector $P(x, y, z)$, velocity V, and orientation O.[5] The scene manager also tracks the listener's pose, typically provided by a head-tracking API.[22]

The Spatial Processing Engine

This module performs the heavy lifting:

1. **Coordinate Transformation:** Translates object positions into a head-relative spherical

coordinate system (θ, ϕ, r).[12]

2. **Occlusion Engine:** Casts rays from the emitter to the listener to check for geometric occluders.[12]
3. **Reflection Solver:** Calculates early reflection paths using the Image Source Method.[16]
4. **Parameter Update:** Interpolates HRTF coefficients and gain values to match the new coordinates.[15]

The Output Rendering Stage

The final stage is the signal processing kernel:

- **Binaural Convolution:** Applies the HRTF filters to each source.[12]
- **Environmental FX:** Adds global late reverberation via an FDN.[18]
- **Limiting and Downmixing:** Sums all virtualized sources into a stereo bus, applying a peak limiter to prevent digital clipping.[26]

Practical Implementation Roadmap

Building a spatial audio system in software is best approached through iterative stages of complexity.

Phase 1: The MVP (Basic Spatial Positioning)

The initial version focuses on azimuthal and elevation positioning for a single mono source using a generic HRTF.

- **Step 1:** Ingest a standard SOFA (Spatially Oriented Format for Acoustics) file from a public database like the SADIE II.[39]
- **Step 2:** Implement a frequency-domain convolution engine using a library like FFTW or the platform's native DSP library (e.g., Apple's vDSP).[39]
- **Step 3:** Implement barycentric interpolation between the nearest HRTF grid points.[13]
- **Outcome:** A mono file can be "panned" anywhere on a 3D sphere around the user's head in real-time.

Phase 2: Intermediate (Depth and Envelopment)

The second phase adds the environmental cues necessary for "out-of-head" localization.

- **Step 1:** Implement a simple "shoebox" room model with adjustable wall absorption coefficients.[16]
- **Step 2:** Integrate a 1st-order Image Source Method to generate six early reflections per source.[16]
- **Step 3:** Add a global FDN reverberator with a T60 decay time matched to the room size.[18]

- **Outcome:** The sound source appears to be at a specific distance within a believable physical room.

Phase 4: Advanced (Personalization and Dynamic Interaction)

The final phase optimizes the system for a wide range of users and devices.

- **Step 1:** Implement SIMD-optimized processing loops to support 64+ simultaneous objects.[58]
- **Step 2:** Integrate a head-tracking data flow using Bluetooth IMU data (e.g., from AirPods or a phone's gyroscope) and apply predictive lag compensation.[62]
- **Step 3:** Add an HRTF individualization module, perhaps by allowing the user to select from a set of "best-fit" HRTF clusters derived from machine learning analysis of ear shapes.[13]
- **Outcome:** A high-performance, personalized spatial audio engine suitable for gaming, VR, or immersive music streaming.

Comparison with Simple Psychoacoustic Effects

It is important to clarify that spatial audio is fundamentally different from traditional "stereo widening" or "surround sound" effects.[4]

- **Stereo Widening:** Often uses phase-shuffling (e.g., adding an out-of-phase L signal to the R channel) to trick the ears into perceiving a wider field. This often causes "hollowing" of the center image and cannot provide height or rear cues.
- **EQ-Based Surround:** Uses static low-pass or high-pass filters to simulate a "rear" channel. This is a crude approximation that lacks the specific timing (ITD) and individualized spectral cues (HRTF) of true spatial audio.
- **True Spatial Audio (Binaural Synthesis):** Reconstructs the entire acoustic wave field as it would appear at the ear canal.[2] It is the only method that can provide precise, stable, and believable 3D localization across all axes, including elevation and depth.[1]

Limitations and Persistent Engineering Challenges

Despite the sophistication of current software engines, several challenges remain:

1. **HRTF Mismatch:** No generic HRTF can perfectly match every user's anatomy.[13] Even with personalization, the "externalization" effect often collapses if the room modeling is not perfectly tuned to the listener's expectations.[43]
2. **Phase Issues:** Interpolating between HRTFs as a source moves rapidly can introduce phase artifacts or "zipper noise".[10] This is often mitigated through crossfading filters, which doubles the CPU load during transitions.[10]
3. **Timbral Coloration:** HRTFs are essentially very complex EQ filters.[13] If the HRTF set has

high "coloration," it can change the perceived timbre of a vocal or instrument, a common complaint among audiophiles regarding spatial audio music mixes.[14]

4. **Battery and Thermal Constraints:** Real-time spatial processing is one of the most CPU-intensive tasks a mobile device can perform.[38] Developers must often trade-off reflection orders or object counts to maintain device thermal stability and battery life.[38]

In conclusion, the successful replication of technologies like Dolby Atmos or Sony 360 Reality Audio in a software-only environment depends on the tight integration of psychoacoustic theory, geometric acoustics, and highly optimized DSP.[5] By utilizing standard HRTF formats like SOFA, implementing low-latency convolution engines, and adding physically-informed room modeling, software developers can provide immersive, 3D auditory experiences to any user with a standard pair of headphones.[3]

Works cited

1. ASAudio: A Survey of Advanced Spatial Audio Research - ACL Anthology, accessed on March 26, 2026, https://aclanthology.org/2025.ijcnlp-long.25.pdf
2. ASAudio: A Survey of Advanced Spatial Audio Research - arXiv, accessed on March 26, 2026, https://arxiv.org/html/2508.10924v2
3. The essential guide to binaural simulation for Dolby Atmos - Audient, accessed on March 26, 2026, https://audient.com/tutorial/the-essential-guide-to-binaural-simulation-for-dolby-atmos/
4. What is spatial audio? How it works and how to use it - Dolby, accessed on March 26, 2026, https://www.dolby.com/experience/home-entertainment/articles/what-is-spatial-audio/
5. 360 Reality Audio - What is the Sony 360RA experience in detail? - VRTONUNG, accessed on March 26, 2026, https://www.vrtonung.de/en/sony-360-reality-audio/
6. Sony 360 Reality Audio | So Immersive. So Real., accessed on March 26, 2026, https://electronics.sony.com/360-reality-audio
7. Sony Introduces All New "360 Reality Audio" Music Experience That Immerses Listeners in a Three-Dimensional Sound Field Powered by Object-Based Spatial Audio Technology, accessed on March 26, 2026, https://www.sony.com/en/SonyInfo/News/Press/201901/19-002E/
8. Enhancing Audio Experiences with Binaural Rendering: How MPEG-H 3D Audio Transforms Soundscapes - Homestyler, accessed on March 26, 2026, https://www.homestyler.com/article/enhancing-audio-experiences-with-binaural-rendering
9. Spatial Sound Mapping: The Invisible Architecture of Immersive Audio - inairspace, accessed on March 26, 2026, https://inairspace.com/blogs/learn-with-inair/spatial-sound-mapping-the-invisible-architecture-of-immersive-audio

10. Tutorial — visr documentation - at www.cvssp.org., accessed on March 26, 2026, https://cvssp.org/data/s3a/public/VISR/visr_installers/macosx/build_py36/doc/userdoc/html/binaural-synthesis-toolkit.html
11. Spatial Sound Rendering for Audio Augmented Reality - IntechOpen, accessed on March 26, 2026, https://www.intechopen.com/chapters/1219737
12. Programmer's Guide — Steam Audio C API documentation, accessed on March 26, 2026, https://valvesoftware.github.io/steam-audio/doc/capi/guide.html
13. HRTFformer: A Spatially-Aware Transformer for Personalized HRTF Upsampling in Immersive Audio Rendering - arXiv, accessed on March 26, 2026, https://arxiv.org/html/2510.01891v1
14. Reverse Engineering Dolby Atmos Binaural Rendering Reveals ..., accessed on March 26, 2026, https://www.grathwohl.me/dolby-atmos-binaural-paper.pdf
15. US11122384B2 - Devices and methods for binaural spatial processing and projection of audio signals - Google Patents, accessed on March 26, 2026, https://patents.google.com/patent/US11122384B2/en
16. Room Impulse Response Simulation with the Image-Source Method and HRTF Interpolation - MATLAB & Simulink - MathWorks, accessed on March 26, 2026, https://www.mathworks.com/help/audio/ug/room-impulse-response-simulation-with-image-source-method-and-hrtf-interpolation.html
17. Wayverb - Image-source Model, accessed on March 26, 2026, https://reuk.github.io/wayverb/image_source.html
18. Feedback Delay Network Architecture - Emergent Mind, accessed on March 26, 2026, https://www.emergentmind.com/topics/feedback-delay-network-fdn-architecture
19. Image Source Approach to Dynamic Early Reflections | Audiokinetic Blog, accessed on March 26, 2026, https://www.audiokinetic.com/rendertron/render/https://www.audiokinetic.com/image-source-approach-to-dynamic-early-reflections/
20. SpatialAudioComponent | Apple Developer Documentation, accessed on March 26, 2026, https://developer.apple.com/documentation/RealityKit/SpatialAudioComponent
21. Dolby Atmos Audio Best Practices - UMG Content Guide, accessed on March 26, 2026, https://contentguide.universalmusic.com/dolby-atmos-audio-best-practices/
22. Personalizing spatial audio in your app | Apple Developer ..., accessed on March 26, 2026, https://developer.apple.com/documentation/phase/personalizing-spatial-audio-in-your-app
23. Dolby Atmos Documentation, accessed on March 26, 2026, https://professional.dolby.com/gaming/gaming-getting-started/dolby-atmos-documentation/
24. Learn - MPEG-H Audio, accessed on March 26, 2026, https://mpegh.com/academy/learn/
25. MPEG-H 3D Audio - Wikipedia, accessed on March 26, 2026,

https://en.wikipedia.org/wiki/MPEG-H_3D_Audio
26. IMMERSIVE AUDIO WITH MPEG 3D AUDIO – STATUS AND OUTLOOK, accessed on March 26, 2026, https://www.iis.fraunhofer.de/content/dam/iis/en/doc/ame/Conference-Paper/NeuendorfM_NAB2014_MPEGH3DAudio.pdf
27. Development of the MPEG-H TV Audio System for ATSC 3.0, accessed on March 26, 2026, https://www.iis.fraunhofer.de/content/dam/iis/en/doc/ame/Conference-Paper/BleidtR-IEEE-2017-Development-of-MPEG-H-TV-Audio-System-for-ATSC-3-0.pdf
28. Quality Enhancement of MPEG-H 3DA Binaural Rendering Using a ..., accessed on March 26, 2026, https://www.mdpi.com/2079-9292/11/9/1491
29. Spatial Audio - Apple Support, accessed on March 26, 2026, https://support.apple.com/en-asia/guide/pro-display-xdr/apd0d6ee8bae/mac
30. Control Spatial Audio and head tracking - Apple Support, accessed on March 26, 2026, https://support.apple.com/guide/airpods/control-spatial-audio-and-head-tracking-dev00eb7e0a3/web
31. AirPods Pro 3 - Technical Specifications - Apple, accessed on March 26, 2026, https://www.apple.com/airpods-pro/specs/
32. FRAUNHOFER EXPLAINS: MPEG-H IMMERSIVE SOUND FOR BROADCAST, STREAMING, AND MUSIC, accessed on March 26, 2026, https://www.iis.fraunhofer.de/content/dam/iis/de/doc/ame/AES-LA-Webinar-April2020-release.pdf
33. PHASE | Apple Developer Documentation, accessed on March 26, 2026, https://developer.apple.com/documentation/phase/
34. Audio and music | Apple Developer Documentation, accessed on March 26, 2026, https://developer.apple.com/documentation/technologyoverviews/audio-and-music
35. Control Spatial Audio and head tracking on Apple devices, accessed on March 26, 2026, https://support.apple.com/guide/beats/control-spatial-audio-head-tracking-apple-dev7db73ea55/web
36. 360 Reality Audio | So immersive. So real. - Sony Asia Pacific, accessed on March 26, 2026, https://www.sony-asia.com/electronics/360-reality-audio
37. Apple does not secretly analyze sine waves to infer head motion. Instead, airpod... | Hacker News, accessed on March 26, 2026, https://news.ycombinator.com/item?id=45403700
38. upHear Mobile Audio - Fraunhofer-Institut für Integrierte Schaltungen IIS, accessed on March 26, 2026, https://www.iis.fraunhofer.de/en/ff/amm/consumer-electronics/uphear/uphear-mobile-audio.html
39. interpolateHRTF - 3-D head-related transfer function (HRTF) interpolation - MATLAB, accessed on March 26, 2026, https://www.mathworks.com/help/audio/ref/interpolatehrtf.html
40. Read, Analyze, and Process SOFA Files - MATLAB & Simulink - MathWorks,

accessed on March 26, 2026, https://www.mathworks.com/help/audio/ug/read-analyze-and-process-sofa-files.html
41. Block diagram of the 3D binaural sound reproduction system incorporating head tracking and room simulation - ResearchGate, accessed on March 26, 2026, https://www.researchgate.net/figure/Block-diagram-of-the-3D-binaural-sound-reproduction-system-incorporating-head-tracking_fig1_228747180
42. Low Latency Audio Processing - QMRO, accessed on March 26, 2026, https://qmro.qmul.ac.uk/xmlui/bitstream/handle/123456789/44697/WANG_Yonghao_Final_PhD_070918.pdf?sequence=1
43. Usability of Individualized Head-Related Transfer Functions in Virtual Reality: Empirical Study With Perceptual Attributes in Sagittal Plane Sound Localization - PMC, accessed on March 26, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC7509635/
44. Block diagram of a 3D binaural sound reproduction system incorporating head tracking and room simulation - ResearchGate, accessed on March 26, 2026, https://www.researchgate.net/figure/Block-diagram-of-a-3D-binaural-sound-reproduction-system-incorporating-head-tracking-and_fig2_228888484
45. The Image-Source Method | Treble Technologies, accessed on March 26, 2026, https://docs.treble.tech/geometrical-solver/image-source-method
46. (PDF) Feedback Delay Networks in Artificial Reverberation and Reverberation Enhancement - ResearchGate, accessed on March 26, 2026, https://www.researchgate.net/publication/322951473_Feedback_Delay_Networks_in_Artificial_Reverberation_and_Reverberation_Enhancement
47. FDN Reverberation | Physical Audio Signal Processing - DSPRelated.com, accessed on March 26, 2026, https://www.dsprelated.com/freebooks/pasp/FDN_Reverberation.html
48. Spatial up-sampling of HRTF sets using generative adversarial networks: A pilot study, accessed on March 26, 2026, https://www.frontiersin.org/journals/signal-processing/articles/10.3389/frsip.2022.904398/full
49. Generic HRTFs May be Good Enough in Virtual Reality. Improving Source Localization through Cross-Modal Plasticity - Frontiers, accessed on March 26, 2026, https://www.frontiersin.org/journals/neuroscience/articles/10.3389/fnins.2018.00021/full
50. Standardizing HRTF Datasets - Sched, accessed on March 26, 2026, https://static.sched.com/hosted_files/aesny23/63/Standardizing%20HRTF%20Datasets.pdf
51. Impact of HRTF individualisation and head movements in a real/virtual localisation task, accessed on March 26, 2026, https://arxiv.org/html/2510.09161v1
52. Listening Tests with Individual versus Generic Head-Related Transfer Functions in Six-Degrees-of-Freedom Virtual Reality, accessed on March 26, 2026, https://www.db-thueringen.de/servlets/MCRFileNodeServlet/dbt_derivate_00045942/ilm1-2019200492_055-062.pdf

53. Assessing machine learning techniques for HRTF individualisation - Sonicom, accessed on March 26, 2026, https://www.sonicom.eu/assessing-machine-learning-techniques-for-hrtf-individualisation/
54. Deep Learning for Personalized Binaural Audio Reproduction - arXiv, accessed on March 26, 2026, https://arxiv.org/html/2509.00400v1
55. SONICOM HRTF Dataset, accessed on March 26, 2026, https://www.sonicom.eu/tools-and-resources/hrtf-dataset/
56. The Extended SONICOM HRTF Dataset and Spatial Audio Metrics toolbox - arXiv, accessed on March 26, 2026, https://arxiv.org/html/2507.05053v2
57. Steam Audio vs Resonance vs Oculus Audio vs UE4 Audio - Audio - Epic Developer Community Forums, accessed on March 26, 2026, https://forums.unrealengine.com/t/steam-audio-vs-resonance-vs-oculus-audio-vs-ue4-audio/109647
58. Faster Tape Emulation with SIMD - by Jatin Chowdhury - Medium, accessed on March 26, 2026, https://medium.com/codex/faster-tape-emulation-with-simd-49287d7b24cf
59. Single instruction, multiple data - Wikipedia, accessed on March 26, 2026, https://en.wikipedia.org/wiki/Single_instruction,_multiple_data
60. SIMD vectorization in the engine · godotengine godot-proposals · Discussion #4544 · GitHub, accessed on March 26, 2026, https://github.com/godotengine/godot-proposals/discussions/4544
61. Performance Optimization of Signal Processing Algorithms for SIMD Architectures - Diva-portal.org, accessed on March 26, 2026, http://www.diva-portal.org/smash/get/diva2:1138490/FULLTEXT01.pdf
62. Low Latency Head Tracking For Seamless Immersive Experiences – INAIRSPACE, accessed on March 26, 2026, https://inairspace.com/blogs/learn-with-inair/low-latency-head-tracking-for-seamless-immersive-experiences
63. Real-Time Motion Compensation for Dynamic Dental Implant Surgery - MDPI, accessed on March 26, 2026, https://www.mdpi.com/2077-0383/14/18/6429
64. A Smart Binaural Hearing Aid Architecture Based on a Mobile Computing Platform - MDPI, accessed on March 26, 2026, https://www.mdpi.com/2079-9292/8/7/811
65. US9348410B2 - Perception based predictive tracking for head mounted displays - Google Patents, accessed on March 26, 2026, https://patents.google.com/patent/US9348410B2/en
66. IMU prediction - Void Computing, accessed on March 26, 2026, https://voidcomputing.hu/blog/imu-prediction/
67. The Performance of Inertial Measurement Unit Sensors on Various Hardware Platforms for Binaural Head-Tracking Applications - PMC, accessed on March 26, 2026, https://pmc.ncbi.nlm.nih.gov/articles/PMC9862010/
68. Low latency audio | Android game development, accessed on March 26, 2026, https://developer.android.com/games/sdk/oboe/low-latency-audio
69. The RIEC HRTF Dataset, accessed on March 26, 2026, https://www.riec.tohoku.ac.jp/pub/hrtf/hrtf_data.html

70. SOFA: Spatially Oriented Format for Acoustics - Österreichische Akademie der Wissenschaften, accessed on March 26, 2026, https://www.oeaw.ac.at/isf/forschung/hoeren/laufende-projekte/sofa-spatially-oriented-format-for-acoustics
71. AirPods Pro 2 Frequency Response changes with Spatial Audio, other than just stereo, channel matching, directional, polarity, and head tracking. : r/oratory1990 - Reddit, accessed on March 26, 2026, https://www.reddit.com/r/oratory1990/comments/1rl7z4v/airpods_pro_2_frequency_response_changes_with/